

Name –Aasif mohd

Project – Terraform

I won't add the installation part it is pretty easy can be found on youtube

Task 1. Create an EC2 instance using terraform on AWS.

Provider:- A provider is a plugin that helps Terraform to understand where it has to create the infrastructure

## Solution.

- 1.First setup the provider aws
2. write the code for instance-type,id,key etc.
3. Check the implementation using terraform plan command
4. use the command terraform apply to create an instance

```
1 provider "aws" {
2     region = "us-east-1" # Set your desired AWS region
3 }
4
5 resource "aws_instance" "example" {
6     ami           = "ami-0731becbf832f281e" # Specify an appropriate AMI ID
7     instance_type = "t2.micro"
8     subnet_id     = "subnet-04c371912f90c2d4f"
9     key_name      = "d"
10 }
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS bash - PROJECT-ec2-ins

Plan: 0 to add, 0 to change, 1 to destroy.

Do you really want to destroy all resources?  
Terraform will destroy all your managed infrastructure, as shown above.  
There is no undo. Only 'yes' will be accepted to confirm.

Enter a value: yes

aws\_instance.example: Destroying... [id=i-0e8f017bdb37f3f6f]  
aws\_instance.example: Still destroying... [id=i-0e8f017bdb37f3f6f, 00m10s elapsed]  
aws\_instance.example: Still destroying... [id=i-0e8f017bdb37f3f6f, 00m20s elapsed]  
aws\_instance.example: Still destroying... [id=i-0e8f017bdb37f3f6f, 00m30s elapsed]  
aws\_instance.example: Still destroying... [id=i-0e8f017bdb37f3f6f, 00m40s elapsed]  
aws\_instance.example: Still destroying... [id=i-0e8f017bdb37f3f6f, 00m50s elapsed]  
aws\_instance.example: Still destroying... [id=i-0e8f017bdb37f3f6f, 01m00s elapsed]  
aws\_instance.example: Destruction complete after 1m1s

Destroy complete! Resources: 1 destroyed.

instance is created on aws

Instance state = running

Clear filters

< 1 >

<input checked="" type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ...	Elastic IP
<input checked="" type="checkbox"/>		i-0e8f017bdb37f3f6f	Running	t2.micro	Initializing	View alarms +	us-east-1f	ec2-35-172-216-232.co...	35.172.216.232	-

Now to delete this instance we can use the command-terraform destroy

```
1 provider "aws" {
2   region = "us-east-1" # Set your desired AWS region
3 }
4
5 resource "aws_instance" "example" {
6   ami           = "ami-0731becbf832f281e" # Specify an appropriate AMI ID
7   instance_type = "t2.micro"
8   subnet_id     = "subnet-04c371912f90c2d4f"
9   key_name      = "d"
10 }
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS bash - PROJECT-ec2-ins

```
}
}

Plan: 0 to add, 0 to change, 1 to destroy.

Do you really want to destroy all resources?
Terraform will destroy all your managed infrastructure, as shown above.
There is no undo. Only 'yes' will be accepted to confirm.

Enter a value: yes

aws_instance.example: Destroying... [id=i-0e8f017bdb37f3f6f]
aws_instance.example: Still destroying... [id=i-0e8f017bdb37f3f6f, 00m10s elapsed]
aws_instance.example: Still destroying... [id=i-0e8f017bdb37f3f6f, 00m20s elapsed]
aws_instance.example: Still destroying... [id=i-0e8f017bdb37f3f6f, 00m30s elapsed]
aws_instance.example: Still destroying... [id=i-0e8f017bdb37f3f6f, 00m40s elapsed]
aws_instance.example: Still destroying... [id=i-0e8f017bdb37f3f6f, 00m50s elapsed]
aws_instance.example: Still destroying... [id=i-0e8f017bdb37f3f6f, 01m00s elapsed]
aws_instance.example: Destruction complete after 1m1s

Destroy complete! Resources: 1 destroyed.
```

## Task 2----

What is a variable in terraform?

A **variable** in Terraform is like a placeholder. It lets you **store a value** (like a region name, server type, etc.) so you can **reuse** it in your code.

Eg. ----

```
variable "region" {
  default = "us-east-1"
}
```

What is an Output ?

An **output** or **output variable** shows **useful information** after you run terraform apply. It's like Terraform saying:

**“Here’s the result you might want to see!”**

Why Use Output ?

- To see values like IP addresses, bucket names, or resource IDs.
- To pass values between modules.
- To debug and check values.

```
output "bucket_name" {  
  
  value = aws_s3_bucket.my_bucket.bucket  
  
}
```

### **Multicloud in Terraform** (Multiple Cloud Providers)

- **What it means:** Using **more than one cloud provider** in the same Terraform project.
- **Example:** You deploy a web app on **AWS** and a database on **Azure**.
- **Why it's used:** To avoid putting all your services on one provider (vendor lock-in), or to take advantage of features from different providers.

#### **In Terraform:**

You use **multiple provider blocks**.

Example-

```
provider "aws" {  
  
  region = "us-east-1"  
  
}  
  
provider "azurerm" {  
  
  features = {} }
```

### **Multiregion in Terraform** (Multiple Regions in One Cloud)

- **What it means:** Deploying resources in **multiple regions** of the **same cloud provider** (like AWS).
- **Example:** Deploying servers in **us-east-1** and **us-west-1** to improve performance or reliability.
- **Why it's used:** For disaster recovery, performance, or compliance.

#### **In Terraform:**

You use **multiple provider blocks** for the same cloud but with different regions.

Example-

```
provider "aws" {  
  
  alias = "us_east"  
  
  region = "us-east-1"  
  
}  
  
provider "aws" {  
  
  alias = "us_west"  
  
  region = "us-west-1"  
  
}  
  
resource "aws_instance" "east_server" {  
  
  provider = aws.us_east  
  
  ...  
  
}  
  
resource "aws_instance" "west_server" {  
  
  provider = aws.us_west  
  
  ...  
  
}
```

## Module

What is a Module in Terraform?

A **module** in Terraform is a **folder** that contains Terraform code and can be **reused** in different places by passing values to it.

A **module** is just a **folder** that contains Terraform code (like resources, variables, outputs) — it's like a **reusable component**.

Think of it like a **Lego block**:

- You build it once.
- You can plug it into multiple places.
- You can pass values to it, and it gives outputs back.

### Why use Modules?

- To **organize** your code.
- To **reuse** the same setup (e.g., EC2 instance, VPC) in multiple environments.
- To **avoid copy-pasting** the same code everywhere.

**Value from main.tf** → goes into → **variable block in module** → used in → **resource or logic inside module's main.tf**

We can use the module code from the outside of module folder easily without have to write the complex code

modules/

└─ ec2\_instance/

├─ main.tf # Contains resources

├─ variables.tf # Defines input variables

└─ outputs.tf # Defines output values

```

modules/ec2_instance
├── main.tf
├── outputs.tf
├── variables.tf
└── main.tf

├─ main.tf # Contains resources
├─ variables.tf # Defines input variables
└─ outputs.tf # Defines output values

main.tf
1  provider "aws" {
2    region = "us-east-1"
3  }
4
5  module "ec2_instance" {
6    source = "../modules/ec2_instance"
7    ami_value = "ami-053b0d53c279acc90"
8    instance_type_value = "t2.micro"
9    subnet_id_value = "subnet-019ea91ed9b5252e7"
10  }

```

## What is a Terraform State File?

When you use **Terraform** to create cloud resources (like servers or databases), it needs to **remember what it created**. That's what the **state file** does!

- It's a file called `terraform.tfstate`.
- It keeps track of **all the resources** Terraform manages.
- Without it, Terraform wouldn't know what exists or what needs to be changed.

### 1. Security Risk (Sensitive Data)

- The state file may contain **secrets** like passwords, API keys, or cloud resource IDs.
- If pushed to GitHub, even in private repos, there's a risk of **accidental leaks**.

--A developer might **clone** the repo to another machine or **copy/paste** it elsewhere (like in a public issue, gist, or Slack).

### 2. No Locking or Sync

- GitHub doesn't support **locking** — if two people push or pull at the same time, it can cause **conflicts or corrupted state**.
- It's not designed for **real-time collaboration** on infrastructure.

## How S3 Fixes These Issues

### 1. Secure & Encrypted Storage

- S3 allows you to enable **server-side encryption**.
- You can set **IAM permissions** to tightly control who can access the file.

### 2. Supports Locking & Versioning

- With **S3 + DynamoDB**, Terraform can **lock the state** so only one person can change it at a time.
- **Versioning** in S3 means you can roll back to older state files if something breaks.

## Workflow: Using GitHub + S3 for Terraform

### 1. Terraform Code in GitHub

- a. The DevOps engineer pulls only the **Terraform code** (like .tf files) from GitHub.
- b. The **state file is not in GitHub** — it's safely stored in **S3**.

### 2. Run terraform init

- a. This connects Terraform to the **remote backend (S3)** using the config in backend "s3" block.
- b. Terraform now knows where the state lives.

### 3. Make Changes & Run terraform apply

- a. Terraform compares the **current state in S3** with your updated code.
- b. It updates infrastructure and automatically **writes the new state** back to S3.

### 4. State File is Updated in S3 Automatically

- a. You don't need to download or upload the state manually.
- b. S3 handles storage, and **DynamoDB (optional)** handles locking.

### 5. If You Make a Mistake? No Worries!

- a. If something breaks, you can **restore a previous version** of the state file from S3's **versioning** feature.
- b. That makes it easy to recover from mistakes.

**if you're not using remote state** (like S3) and someone forgets to **push or update the state file**, there's a **real risk of duplicate resources being created**.

Let's break it down:

## When Using Local State and GitHub Only (Bad Practice)

### 1. Engineer A

- a. Clones the repo.
- b. Applies Terraform with **local state file**.
- c. Creates, for example, an EC2 instance.
- d. **Forgets to push the updated state file** to GitHub (or it's not in GitHub at all).

### 2. Engineer B

- a. Pulls the same repo, **but without the latest state**.
- b. Terraform sees **no record** of the EC2 instance.

- c. So when B runs terraform apply, it **creates another EC2 instance** — a **duplicate**.

What does "provision" mean in Terraform?

In **Terraform**, **provisioning** means **setting up or configuring** a resource **after** it's created. It's like saying:

“Hey Terraform, after you create this server, run this script to install software or make some changes inside it.”

Example to understand it:

Imagine you're using Terraform to create a **virtual machine (VM)** in the cloud. That's easy:

```
resource "aws_instance" "example" {  
  ami      = "ami-123456"  
  instance_type = "t2.micro"  
}
```

This creates a VM. But what if you want that VM to:

- Install **Nginx**?
- Set up **some files**?
- Run a **shell script**?

That's where **provisioners** come in.

✂ Example with a provisioner:

```
resource "aws_instance" "example" {  
  ami      = "ami-123456"  
  instance_type = "t2.micro"
```



```
provisioner "remote-exec" {  
  
  inline = [  
  
    "sudo apt update",  
  
    "sudo apt install -y nginx"  
  
  ]  
  
}  
  
}
```

This tells Terraform:

“After you create the VM, connect to it and run these commands.”

#### Types of Provisioners:

1. **file** – Uploads a file to the server
2. **remote-exec** – Runs commands on the server (SSH)
3. **local-exec** – Runs a command on your local machine (not the server)

#### ◇ Terraform Workspaces (Quick Explanation)

**Terraform workspaces** let you use the **same code** for different environments (like dev, stage, prod) by creating a **separate state file** for each one.

So:

- One codebase ✓
- Multiple environments ✓
- Separate .tfstate files ✓

#### Why Use Workspaces?

You might have one Terraform codebase (like for setting up a server), but want:

- One for **development** (dev)

- One for **production** (prod)

Instead of copying the code multiple times, you can:

- Keep **one codebase**
- Use **different workspaces** for each environment

Each workspace has its **own separate state**.

Command	What It Does
terraform workspace list	Shows all available workspaces
terraform workspace new dev	Creates a new workspace named dev
terraform workspace select dev	Switches to the dev workspace
terraform workspace show	Shows the current workspace

## Terraform Secrets with Vault

**Vault** is a tool to **securely store secrets** like passwords, API keys, and access tokens.

Instead of putting secrets directly in your Terraform code (which is risky), you can:

1. **Store secrets in Vault**
2. **Configure Terraform to read those secrets**
3. **Use them in your infrastructure setup**

## Step-by-Step Process

1. **Secrets Stored in Vault**
  - a. Example: Store db\_password = "supersecret" in Vault at path like secret/data/db.
2. **Terraform Authenticates with Vault**
  - a. Terraform uses an **AppRole, Token**, or other auth method.
  - b. Vault checks if Terraform is allowed to access the secret.
3. **Terraform Reads Secret**
  - a. Terraform uses the **Vault provider** to fetch the secret.
  - b. The secret is used during provisioning (like setting up a database).
4. **No Secrets in Code!**

- a. The sensitive data stays in Vault.
- b. Terraform just reads it at runtime.

### Why Use Vault with Terraform?

- **Security:** No hardcoding secrets in code
- **Central management:** Keep all secrets in one place
- **Auditing:** See who accessed what and when